

Quick and Tidy Guide to Econometrics

Ryan Safner

Fall 2019

Contents

Overview	2
Summary Cheatsheet	3
General Tips	3
Data Import	3
Data Wrangling	3
Regressions	3
Plotting	3
R Basics	5
Data Wrangling	6
Import	6
Looking at Data	6
General Data Manipulation	7
Subset Data	8
Transform Data	9
Summarize Data	9
Grouped-Summaries	10
Categorical Data	10
Correlation	11
Regression	13
Interpretation of Output	14
Tidying Output with Broom	14
Extensions	15
Categorical Data: Dummy Variables	15
Interaction Terms	16
Polynomial Models	17
Logarithmic Models	17
Regression Output Tables	18
Plots	20
Histogram	20
Scatterplot (with Regression Line)	21
Customized Example	21

Overview

This is a quick guide to using R for basic econometrics and data analysis tasks (i.e. manipulating data, running regressions, and making plots and tables) in an *opinionated* way, using the `tidyverse` packages and grammar. Use this as document a shortcut or cheatsheet to refer to the packages, commands, and syntax necessary to perform these tasks.

It is wise to always **load `tidyverse`**¹ at the beginning of each R session or markdown document, as most of the guide makes use of packages and commands that are *assumed* to already be loaded with `tidyverse`.

```
library(tidyverse)
```

Throughout this guide, I use an `example_data.csv` file for data. To see how I made this data with R, see the Appendix.

¹With `library(tidyverse)`.

Summary Cheatsheet

General Tips

- Make an R Project to organize files
- Always start with `library(tidyverse)`
- Never `View()` or `install.packages()` in a .Rmd file

Data Import

```
df<-read_*("file.*")
```

Data Wrangling

Command	Does	Example
<code>select()</code>	Keep desired columns (variables)	<code>gapminder %>% select(pop)</code>
<code>filter()</code>	Keep desired rows (observations)	<code>gapminder %>% filter(country="France")</code>
<code>arrange()</code>	Reorder rows (e.g. in numerical order)	<code>gapminder %>% arrange(pop)</code>
<code>mutate()</code>	Create new variables	<code>gapminder %>% mutate(GDP = gdpPerCap * pop)</code>
<code>summarize()</code>	Collapse data into summary statistics	<code>gapminder %>% summarize(avg_GDP = mean(gdpPerCap))</code>
<code>group_by()</code>	Perform any of the above functions by groups/categories	<code>gapminder %>% group_by(country) %>% summarize(avg_LE = mean(LifeExp))</code>

Regressions

```
reg<-lm(y~x+z, data = df)
```

- Extensions:
 - Polynomial: `I(var_name^2)`
 - Log: `log(var_name)`
 - Dummy/Fixed Effects (if not already a factor or dummy): `factor(category_var)`
 - Interaction: `var_1:var_2`
- Viewing:
 - `summary(reg)` for full output
 - `broom::tidy(reg)` to view coefficients in tidy tibble
 - `broom::glance(reg)` to view regression statistics in tidy tibble
 - `broom::augment(reg)` to add regression-based observations (i.e. \hat{u}_i, \hat{y}_i)
- Making Regression Tables:
 - `huxtable::huxreg(reg)`

Plotting

```
ggplot(data = df)+ # layer defining dataframe for data source
# aesthetics layer to map variables to aesthetics
```

```
aes(x = X,  
    y = Y,  
    color = shape)+ # will color by shape  
# geometries layer(s)  
geom_point()+  
geom_smooth(method = "lm") # add a regression line
```

R Basics

- R is “object-oriented”:
 - Assign values in objects: `my_object ← my_values`
 - Overwrite objects: `my_object ← my_new_values`
 - Run functions on objects: `function_name(my_object)`
- Data types:
 - numeric (“double” or “integer”) data are numbers
 - * can use for math and statistics
 - character: strings of text
 - * values must always be “in quotes”
 - factor: indicates membership in one of several possible categories or groups
- Object types:
 - vector: collection of objects
 - * create with `c()` function
 - data.frame or tibble: each row is a vector of same data type²
 - * rows are observations
 - * columns are variables
- Packages:
 - Install any package with `install.packages("package_name")`
 - * Only necessary once, if package doesn't already exist
 - Load a package for each session needed with `library("package_name")`
 - Description of packages we use:

Package	Use
tidyverse	For tibbles (tibble), %>% operator (magrittr), data import (readr), data wrangling (dplyr), plotting (ggplot2)
broom	For tidy regression outputs
huxtable	For making regression tables
car	For regression tests (heteroskedasticity, outliers, F-test)
lmtest	
estimatr	For regression with heteroskedasticity-robust standard errors

²Henceforth, with tidyverse, I refer to all data.frames as tibbles

Data Wrangling

Import

- Import data with `read_*()` where the `*` represents the file extension (e.g. `csv`, `tsv`, `xls`, `xlsx`, `dta`)³ and inside the parentheses you place the location of the file on your computer or web URL in quotes
- Be sure to assign your data to a tibble!

```
#my_df<-read_csv("https://metricsf19.classes.ryansafner.com/data/example_data.csv")

my_df<-read_csv("example_data.csv")
```

Looking at Data

- type the name of the tibble to print its contents
- `str()` gives the structure of a tibble
- `head()` prints the first few rows of a tibble
- `View()` will open the tibble in a separate window for inspection⁴
- `glimpse()` gives the structure in a horizontal way

```
my_df
```

```
## # A tibble: 100 x 5
##       X       Z       U Shape      Y
##   <dbl> <dbl> <dbl> <chr>   <dbl>
## 1  9.74  10.8  0.576 Square  9.60
## 2 11.8   12.0  0.991 Circle  2.58
## 3 10.3   10.7 -0.781 Circle  4.91
## 4  9.03  14.4  0.412 Triangle 24.7
## 5  9.37  18.0 -0.675 Square  34.5
## 6 10.3   14.7  1.23  Square  21.4
## 7  9.57  17.4 -0.0248 Square  32.5
## 8 10.7   16.7 -0.750 Circle  24.6
## 9  9.71  15.5 -0.541 Square  25.0
## 10 10.0   15.0 -0.0772 Square  22.3
## # ... with 90 more rows
```

```
str(my_df)
```

```
## Classes 'spec_tbl_df', 'tbl_df', 'tbl' and 'data.frame': 100 obs. of  5 variables:
## $ X      : num  9.74 11.79 10.33 9.03 9.37 ...
## $ Z      : num  10.8 12 10.7 14.4 18 ...
## $ U      : num  0.576 0.991 -0.781 0.412 -0.675 ...
## $ Shape: chr  "Square" "Circle" "Circle" "Triangle" ...
## $ Y      : num  9.6 2.58 4.91 24.67 34.45 ...
## - attr(*, "spec")=
## .. cols(
## ..   X = col_double(),
## ..   Z = col_double(),
## ..   U = col_double(),
## ..   Shape = col_character(),
```

³Loading tidyverse automatically loads the readr package, allowing you to load csv and tsv files without loading any package. Other file types require loading other packages, such as readxl (for Excel files) or haven or foreign (which are pretty good at reading any other type).

⁴Do not run this command in a markdown document or it will not knit!

```
## .. Y = col_double()
## .. )
```

```
head(my_df)
```

```
## # A tibble: 6 x 5
##       X       Z       U Shape      Y
##   <dbl> <dbl> <dbl> <chr>   <dbl>
## 1  9.74  10.8  0.576 Square    9.60
## 2 11.8   12.0  0.991 Circle    2.58
## 3 10.3   10.7 -0.781 Circle    4.91
## 4  9.03  14.4  0.412 Triangle 24.7
## 5  9.37  18.0 -0.675 Square   34.5
## 6 10.3   14.7  1.23  Square   21.4
```

```
glimpse(my_df)
```

```
## Observations: 100
## Variables: 5
## $ X      <dbl> 9.737951, 11.785850, 10.334175, 9.028176, 9.371203, 10.272492 ...
## $ Z      <dbl> 10.76144, 12.02853, 10.72002, 14.41761, 18.03719, 14.68002, 1 ...
## $ U      <dbl> 0.57631430, 0.99107639, -0.78134698, 0.41238420, -0.67469748, ...
## $ Shape  <chr> "Square", "Circle", "Circle", "Triangle", "Square", "Square", ...
## $ Y      <dbl> 9.598392, 2.579783, 4.905085, 24.673541, 34.452718, 21.389187 ...
```

General Data Manipulation

The following table provides the major verbs for manipulating data. Further sections below provide examples for i. *subsetting* data ii. *transforming* data iii. *summarizing* data

Verb (from dplyr)	Action
<code>select()</code>	Keep desired columns (variables)
<code>filter()</code>	Keep desired rows (observations)
<code>arrange()</code>	Reorder rows (e.g. in numerical order)
<code>mutate()</code>	Create new variables
<code>summarize()</code>	Collapse data into summary statistics
<code>group_by()</code>	Perform any of the above functions by groups/categories

Most verbs allow you to manipulate data according to some **condition(s)**. Popular operators for performing conditional operations are listed in the table below:

Command	Effect
<code><; ></code>	Less than; greater than
<code>≤; ≥</code>	Less than or equal to; greater than or equal to
<code>==; ≠</code>	Is equal to; is not equal to
<code>%in%</code>	Is in the set of [a vector of options]
<code>is.na()</code>	Is missing (NA)

Finally, for each command, you can alternatively:

```
# Just view the output
my_df %>% verb()
```

```
# Assign to a different object
my_df_2<-my_df %>% verb()
my_df_2 # then view the output

# Assign to the original object (and overwrite it)
my_df<-my_df %>% verb()
my_df # then view the output
```

Subset Data

- To *subset* data and take only a portion of the data set by some condition(s) for various purposes, use
 - `select()` to subset by columns (variables)
 - `filter()` to subset by rows (observations)

```
# look only at data for "circles" AND where X>10
my_df %>%
  filter(Shape=="Circle",
         X>10)
```

```
## # A tibble: 17 x 5
##       X      Z      U Shape      Y
##   <dbl> <dbl> <dbl> <chr>   <dbl>
## 1  11.8  12.0  0.991 Circle  2.58
## 2  10.3  10.7 -0.781 Circle  4.91
## 3  10.7  16.7 -0.750 Circle 24.6
## 4  11.5  11.4  0.0275 Circle  1.12
## 5  11.9  12.8  1.02   Circle  4.65
## 6  10.2  15.9  0.0326 Circle 24.7
## 7  11.6  14.3 -1.05   Circle 10.6
## 8  10.4  16.9  0.929 Circle 28.4
## 9  10.6  10.3 -2.05   Circle  0.707
## 10 10.4  13.2 -1.59   Circle 12.3
## 11 11.5  19.2 -0.372 Circle 30.9
## 12 10.9  19.5 -1.54   Circle 33.7
## 13 10.2  15.0 -1.38   Circle 20.3
## 14 11.1  16.7  0.803 Circle 24.4
## 15 10.2  19.6  0.967 Circle 39.0
## 16 11.4  12.4 -0.745 Circle  4.36
## 17 13.0  10.9 -0.232 Circle -12.4
```

```
# look only at X and Y
my_df %>%
  select(X,Y)
```

```
## # A tibble: 100 x 2
##       X      Y
##   <dbl> <dbl>
## 1  9.74  9.60
## 2 11.8   2.58
## 3 10.3   4.91
## 4  9.03 24.7
## 5  9.37 34.5
## 6 10.3  21.4
## 7  9.57 32.5
```



```
## 8 10.7 24.6
## 9 9.71 25.0
## 10 10.0 22.3
## # ... with 90 more rows
```

Transform Data

- Data transformation uses the `mutate()` command to either
 - create a new variable `mutate(new_name = conditions on existing variables)`
 - change a variable (and overwrite it) `mutate(existing_variable = conditions on existing_variable)`

```
# take the log of Y,
# and the log of X,
# and make new variable V, which is 0.5 times X times Z,
# and change the class of Shape from to a character variable to a factor variable
```

```
my_df<-my_df %>%
  mutate(log_Y=log(Y),
         log_X=log(X),
         V = 0.5*(X*Z),
         Shape = as.factor(Shape))
my_df
```

```
## # A tibble: 100 x 8
##       X      Z      U Shape      Y log_Y log_X      V
##   <dbl> <dbl> <dbl> <fct>   <dbl> <dbl> <dbl> <dbl>
## 1  9.74  10.8  0.576 Square   9.60  2.26  2.28  52.4
## 2 11.8   12.0  0.991 Circle   2.58  0.948  2.47  70.9
## 3 10.3   10.7 -0.781 Circle   4.91  1.59  2.34  55.4
## 4  9.03  14.4  0.412 Triangle 24.7   3.21  2.20  65.1
## 5  9.37  18.0 -0.675 Square  34.5   3.54  2.24  84.5
## 6 10.3   14.7  1.23   Square  21.4   3.06  2.33  75.4
## 7  9.57  17.4 -0.0248 Square  32.5   3.48  2.26  83.4
## 8 10.7   16.7 -0.750 Circle  24.6   3.20  2.37  89.7
## 9  9.71  15.5 -0.541 Square  25.0   3.22  2.27  75.4
## 10 10.0   15.0 -0.0772 Square  22.3   3.11  2.30  75.2
## # ... with 90 more rows
```

Summarize Data

- Create summary statistics for datasets with `summarize()`. This will create a tibble of summary statistics.

Below is a table of popular statistics-based commands for summarizing data. Except for the first two, place a variable inside each command, and optionally set it equal to a name for the statistic to be outputted.

Command	Does
<code>n()</code>	Number of observations (nothing goes in parentheses!)
<code>n_distinct()</code>	Number of unique observations (nothing goes in parentheses!)
<code>sum()</code>	Sum all observations of a variable
<code>mean()</code>	Average of all observations of a variable
<code>median()</code>	50th percentile of all observations of a variable

Command	Does
sd()	Standard deviation of all observations of a variable
min()	Minimum value of a variable
max()	Maximum value of a variable
quantile(x, 0.25)	Specified percentile (example 25th percentile) of a variable
first()	First value of a variable
last()	Last value of a variable
nth(x, 2)	Specified position of a variable (example 2nd)

```
# find number of obs, and mean and sd of X and Y
my_df %>%
  summarize(n(),
            Mean_X = mean(X),
            Std_dev_X = sd(X),
            Mean_Y = mean(Y),
            Std_dev_Y = sd(Y))
```

```
## # A tibble: 1 x 5
##   `n()` Mean_X Std_dev_X Mean_Y Std_dev_Y
##   <int> <dbl>   <dbl> <dbl>   <dbl>
## 1   100  10.2     1.05  21.5    12.4
```

Grouped-Summaries

- You can run summary statistics by group by first using `group_by(categorical_variable)` and then `summarize()`:

```
# get mean of X and Y for each Shape
my_df %>%
  group_by(Shape) %>%
  summarize(mean_X = mean(X),
            mean_Y = mean(Y))
```

```
## # A tibble: 3 x 3
##   Shape mean_X mean_Y
##   <fct>   <dbl> <dbl>
## 1 Circle  10.4    21.1
## 2 Square   9.96    21.3
## 3 Triangle 10.3    22.4
```

Categorical Data

- For categorical data (factors), you can quickly produce a frequency table of each category with `count(factor_variable_name)`
- `distinct()` shows the distinct values of a specified variable (often useful for finding the different categories)

```
# count by shape
my_df %>%
  count(Shape)
```

```
## # A tibble: 3 x 2
##   Shape      n
##   <fct> <int>
## 1 Circle    27
```

```
## 2 Square      47
## 3 Triangle    26
# get the distinct shapes
my_df %>%
  distinct(Shape)

## # A tibble: 3 x 1
##   Shape
##   <fct>
## 1 Square
## 2 Circle
## 3 Triangle
```

Correlation

- You can quickly produce a correlation table (2+ variables) so long as they are numeric (i.e. not character or factor):

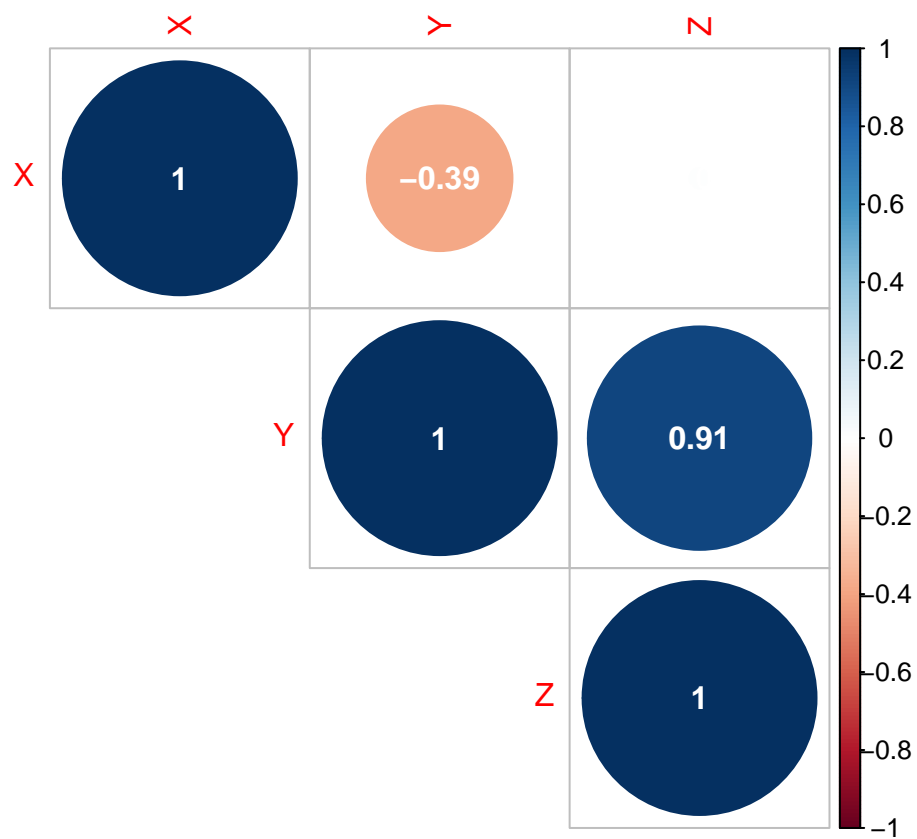
```
my_df %>%
  select(X,Y,Z) %>%
  cor()
```

```
##           X           Y           Z
## X  1.000000000 -0.3872977  0.007279217
## Y -0.387297726  1.0000000  0.912385345
## Z  0.007279217  0.9123853  1.000000000
```

- A nice **correlogram** can be made with the **corrplot** package:

```
library(corrplot)

my_df %>%
  select(X,Y,Z) %>%
  cor() %>%
  corrplot(.,
    method = "circle", # choose circle, square, ellipse, number, pie, shade, color
    type = "upper", # show only upper triangle of matrix
    addCoef.col="white") # add white numbers
```



Regression

The primary task in econometrics is running a regression on data. Regression uses the linear model `lm()` command where “Y” is regressed on all X variables, connected with +s, where the data is sourced from your tibble.

```
lm(Y~X+Z, data = my_df)
```

This will output the coefficients only. To get full information on coefficients, standard errors, hypothesis testing, and regression fit, pipe into the `summary()` command, or save the regression as an object and then run `summary()` on it. I show both methods below:

```
my_df %>%
  lm(data = ., # pipes my_data into the data = argument
    Y~X+Z) %>%
  summary()

##
## Call:
## lm(formula = Y ~ X + Z, data = .)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.2913 -0.8377 -0.0387  0.7535  3.7508
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 13.51488    1.51898   8.897 3.22e-14 ***
## X            -4.65217    0.13328  -34.904 < 2e-16 ***
## Z             3.65230    0.04504   81.089 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.397 on 97 degrees of freedom
## Multiple R-squared:  0.9876, Adjusted R-squared:  0.9874
## F-statistic: 3877 on 2 and 97 DF, p-value: < 2.2e-16

my_reg_1<-lm(Y~X+Z, data = my_df)
summary(my_reg_1)

##
## Call:
## lm(formula = Y ~ X + Z, data = my_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.2913 -0.8377 -0.0387  0.7535  3.7508
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 13.51488    1.51898   8.897 3.22e-14 ***
## X            -4.65217    0.13328  -34.904 < 2e-16 ***
## Z             3.65230    0.04504   81.089 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.397 on 97 degrees of freedom
## Multiple R-squared:  0.9876, Adjusted R-squared:  0.9874
```

```
## F-statistic: 3877 on 2 and 97 DF, p-value: < 2.2e-16
```

Interpretation of Output

- The top row Residuals describes the distribution of the residuals.
- The Coefficients table describes the OLS parameters $\hat{\beta}_j$'s, where each row is a right-hand side variable, starting with (Intercept) ($\hat{\beta}_0$), then ($\hat{\beta}_1$), etc.
 - Estimate column describes the OLS parameters ($\hat{\beta}_0, \hat{\beta}_1, \dots$)
 - Std. Error column describes the standard error of each parameter
 - t value column describes the test statistic for a hypothesis test where H_0 : that particular $\hat{\beta}_j = 0$
 - Pr(>|t|) column is the p-value on that hypothesis test for that parameter (roughly, we're looking for it to be less than 0.05, if it is, there will be * stars to the right of it.
- The bottom three rows describes the goodness of fit of the regression
 - Residual standard error is the Standard Error of the Regression σ_u
 - Multiple R-squared is R^2 , Adjusted R-squared is \bar{R}^2
 - F-statistic is the F-statistic on an All F-test (all betas are equal to 0), and associated p-value on that test

Tidying Output with Broom

- The broom package allows us to output regressions into tidy tibbles that we can easily print, work with, and extract individual parameters or statistics from for further analysis

Command	Does
tidy()	Takes the saved reg_object and makes a tibble of the coefficients table only
augment()	Create dataset with calculated values (e.g. .fitted, .resid)
glance()	Get statistics of regression fit (e.g. r.squared, sigma)

```
# load broom
library(broom)

# tidy to get coefficients in a tidy tibble
my_reg_1_tidy <- tidy(my_reg_1)
my_reg_1_tidy

## # A tibble: 3 x 5
##   term      estimate std.error statistic  p.value
##   <chr>      <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)  13.5      1.52      8.90 3.22e-14
## 2 X          -4.65     0.133    -34.9 1.02e-56
## 3 Z           3.65     0.0450     81.1 6.18e-91

# glance (original lm object) to view statistics
glance(my_reg_1)

## # A tibble: 1 x 11
##   r.squared adj.r.squared sigma statistic  p.value    df logLik   AIC   BIC
##   <dbl>      <dbl> <dbl>    <dbl>    <dbl> <int> <dbl> <dbl> <dbl>
## 1   0.988      0.987  1.40    3877. 2.87e-93     3  -174.  356.  366.
## # ... with 2 more variables: deviance <dbl>, df.residual <int>
```

```
# "r.squared" and "adj.r.squared" are self-explanatory
# "sigma" is the Standard Error of the Regression (SER)
# "statistic" is the F-statistic on the All-F test
# "p.value" is the p-value from that All-F test
```

```
my_reg_1_aug<-augment(my_reg_1)
my_reg_1_aug
```

```
## # A tibble: 100 x 10
##       Y      X      Z .fitted .se.fit .resid  .hat .sigma  .cooksd .std.resid
##   <dbl> <dbl> <dbl>   <dbl>   <dbl>   <dbl>  <dbl> <dbl>    <dbl>    <dbl>
## 1  9.60  9.74  10.8    7.52    0.248    2.08   0.0315  1.39  0.0249     1.51
## 2  2.58 11.8   12.0    2.62    0.293   -0.0372 0.0441  1.40  0.0000114 -0.0272
## 3  4.91 10.3   10.7    4.59    0.244    0.314   0.0306  1.40  0.000547     0.228
## 4 24.7   9.03  14.4   24.2    0.209    0.502   0.0223  1.40  0.00100     0.363
## 5 34.5   9.37  18.0   35.8    0.219   -1.34    0.0246  1.40  0.00797    -0.974
## 6 21.4   10.3  14.7   19.3    0.142    2.05    0.0103  1.39  0.00755     1.47
## 7 32.5   9.57  17.4   32.7    0.192   -0.203   0.0188  1.40  0.000138    -0.147
## 8 24.6   10.7  16.7   24.6    0.174    0.0518  0.0155  1.40  0.00000734  0.0374
## 9 25.0   9.71  15.5   25.1    0.154   -0.0983  0.0121  1.40  0.0000204    -0.0708
## 10 22.3  10.0   15.0   21.7    0.141    0.652   0.0102  1.40  0.000758     0.469
## # ... with 90 more rows
```

```
# ".fitted" are predicted (Y-hat) values from the model
# ".resid" are the residuals (u-hat) for each X-value
```

- `augment` is particularly useful for plotting fitted or residual values, as in a residual plot, where you can set `aes(y = .resid)`.

Extensions

Categorical Data: Dummy Variables

For categorical data, you can run a regression with a dummy variable if that variable takes on the values of 0 or 1.

If the variable has multiple possible categories, you can use (or make) a dummy variable for *each* of the n categories and include all $n - 1$ dummy variables in the regression (to avoid the dummy variable trap!).

If your variable exists as a factor variable (e.g. the value of each observation for that variable is the name of the category), you can simply add that variable in the regression and R will automatically create a dummy for each category and include $n - 1$ dummies in a regression:

```
# run a regression with Shape, a factor variable
## which has "Circle," "Triangle," and "Square" for categories
```

```
shape_reg<-lm(Y~Shape, data = my_df)
summary(shape_reg)
```

```
##
## Call:
## lm(formula = Y ~ Shape, data = my_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -33.446  -8.783  -0.010   10.234   21.661
##
## Coefficients:
```

```
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)    21.069      2.416   8.720 7.74e-14 ***
## ShapeSquare     0.206      3.032   0.068   0.946
## ShapeTriangle   1.319      3.450   0.382   0.703
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 12.55 on 97 degrees of freedom
## Multiple R-squared:  0.00182, Adjusted R-squared: -0.01876
## F-statistic: 0.08841 on 2 and 97 DF, p-value: 0.9155

# Note R made two dummies:
# # ShapeTriangle for Triangle (0 or 1)
# # ShapeSquare for Square (0 or 1)
# # and left out Circle as the reference category
```

Interpretation:

- $\hat{\beta}_0$ is the average value of Y for the reference category
 - e.g. Circles have an average Y of 19.63
- Each $\hat{\beta}$ is the *difference* between that category and the reference category
 - e.g. Triangle has an average Y that is 4.66 larger than Circle
 - e.g. Square has an average Y that is 0.15 larger than Circle

Interaction Terms

To interact two variables in a regression and create an interaction term, simply add them to the regression with : or * between them.

```
interact_reg<-lm(Y~X+Z+X:Z, data = my_df)
summary(interact_reg)

##
## Call:
## lm(formula = Y ~ X + Z + X:Z, data = my_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.01411 -0.75425 -0.06623  0.81317  2.76908
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  47.28197    5.33968   8.855 4.28e-14 ***
## X           -7.99364    0.52518 -15.221 < 2e-16 ***
## Z            1.28283    0.36586   3.506 0.000693 ***
## X:Z          0.23436    0.03599   6.511 3.41e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.169 on 96 degrees of freedom
## Multiple R-squared:  0.9914, Adjusted R-squared:  0.9912
## F-statistic: 3701 on 3 and 96 DF, p-value: < 2.2e-16
```


Polynomial Models

To run a polynomial regression, simply add a higher order variable, which you can first `mutate()`, or use the `I()` command to create a quadratic (or higher order) term in your regression:

```
reg_quad<-lm(Y~X+I(X^2), data = my_df)
summary(reg_quad)
```

```
##
## Call:
## lm(formula = Y ~ X + I(X^2), data = my_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -20.487  -8.755  -1.634   11.213   19.902
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -120.513      75.032  -1.606   0.1115
## X              32.862      14.779   2.224   0.0285 *
## I(X^2)        -1.839       0.724  -2.540   0.0127 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 11.22 on 97 degrees of freedom
## Multiple R-squared:  0.203, Adjusted R-squared:  0.1866
## F-statistic: 12.35 on 2 and 97 DF, p-value: 1.663e-05
```

Logarithmic Models

To run a regression with a logged variable, you can first `mutate()` the logged variable, or use the `log()` command to create a logged variable in your regression:

```
reg_log_log<-lm(log(Y)~log(X), data = my_df)
summary(reg_log_log)
```

```
##
## Call:
## lm(formula = log(Y) ~ log(X), data = my_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.9873 -0.2975  0.0962  0.6634  1.1164
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  10.7762      1.9425   5.548 2.51e-07 ***
## log(X)        -3.4454      0.8396  -4.103 8.48e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.85 on 97 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.1479, Adjusted R-squared:  0.1391
## F-statistic: 16.84 on 1 and 97 DF, p-value: 8.482e-05
```

Regression Output Tables

There are several methods to make regression output tables, but I have found best use with the `huxtable` package's `huxreg()` command. There are three parts to the command (each separated by commas):

1. Add your regression `lm` objects, separated by commas.
 - Optionally define a "column name" = `your_lm_object` for each.
2. Optionally rename and omit your X-variables as desired inside `coefs = c()`
 - To change a variable's name down each row, set "desired name" = `var_name`.
 - Any X-variable not listed inside `coefs = c()` will be omitted from the table!
3. Optionally rename and omit statistics as desired inside `stats = c()`
 - To change a variable's name down each row, set "desired name" = `stat_name`.
 - Any statistics not listed inside `stats = c()` will be omitted from the table!

```
# default example
library(huxtable)
my_reg_1 %>% huxreg()
```

	(1)
(Intercept)	13.515 *** (1.519)
X	-4.652 *** (0.133)
Z	3.652 *** (0.045)
N	100
R2	0.988
logLik	-173.792
AIC	355.584

*** p < 0.001; ** p < 0.01; * p < 0.05.

```
# heavily customized
library(huxtable)
huxreg(my_reg_1,
  shape_reg,
  interact_reg,
  reg_quad,
  reg_log_log,
  coefs = c("Constant" = "(Intercept)",
    "X" = "X",
    "Z" = "Z",
    "Square" = "ShapeSquare",
    "Triangle" = "ShapeTriangle",
    "X:Z", "X:Z",
    "$X^2$" = "I(X^2)",
    "ln(X)" = "log(X)"),
  statistics = c("N" = "nobs",
    "$R^2$" = "r.squared",
    "SER" = "sigma"),
  note = NULL, # suppress footnote for stars, to insert Fixed Effects Row below
  number_format = 3) %>% # round to three decimals
```

```

# Adding a lot of custom parts to table below

add_rows(c("Fixed Effects", "None", "Shape", "None", "None", "None"), # add fixed effects row
  after = 17) %>% # insert after 17th row

# allow math to render as R^2 and X^2
set_escape_contents(c(14,20), c(1,1), FALSE) %>% # R^2 is rows 14 and 20, column 1

# add centered "Y" in second row
insert_row(c("", rep("Y", 4)), "ln(Y)", after = 1) %>%
merge_cells(c(2,2,2), 2:5) %>%

# create borders
set_all_borders(0) %>% # remove all borders to manually set my own
set_top_border(1, 1:6, 2) %>% # add border size 2 to first row, columns 1:6
set_top_border(2, 1:6, 1) %>% # add border size 1 to second row, columns 1:6
set_top_border(3, 2:5, 1) %>% # add border size 1 to second row, columns 2:5
set_top_border(19, 2:6, 1) %>% # add border size 1 to second row, columns 2:5
set_bottom_border(22, 1:6, 2) %>% # add border size 2 to 22nd row, columns 1:6

# caption the table
set_caption( "Regression Results")

```

Table 7: Regression Results

	(1)	(2)	(3)	(4)	(5)
	Y				ln(Y)
Constant	13.515 *** (1.519)	21.069 *** (2.416)	47.282 *** (5.340)	-120.514 (75.032)	10.776 *** (1.942)
X	-4.652 *** (0.133)		-7.994 *** (0.525)	32.862 * (14.779)	
Z	3.652 *** (0.045)		1.283 *** (0.366)		
Square		0.206 (3.032)			
Triangle		1.319 (3.450)			
X:Z			0.234 *** (0.036)		
X ²				-1.839 * (0.724)	
ln(X)					-3.445 *** (0.840)
Fixed Effects	None	Shape	None	None	None
N	100	100	100	100	99
R ²	0.988	0.002	0.991	0.203	0.148
SER	1.397	12.554	1.169	11.218	0.850


Plots

Plots of any kind can be made with `ggplot2`, which uses a “grammar of graphics” to build plots layer by layer. Some possible layers are described in the table below, *required* layers are boldened:

Command	Layer	Description
<code>data =</code>	Data	Defines what tibble to use for the data
<code>+aes()</code>	Aesthetics	Defines what variables from data will be mapped to markings
<code>+geom_*()</code>	Geometry	Defines what markings to make, e.g. point, histogram, line, smooth (for regression lines)
<code>+coord_*()</code>	Coordinates	Scales for axes
<code>+scale_*()</code>	Scales	Define the range of values
<code>+facet_*()</code>	Facets	Group into subplots


Histogram

```
ggplot(data = my_df)+  
  aes(x = X)+  
  geom_histogram()
```



Scatterplot (with Regression Line)

```
ggplot(data = my_df)+  
  aes(x = X,  
      y = Y)+  
  geom_point()+  
  geom_smooth(method = "lm")
```



Customized Example

```
ggplot(data = my_df)+  
  aes(x = X,  
      y = Y)+  
  geom_point(aes(shape = Shape,  
                color = Shape),  
            size = 2)+  
  geom_smooth(method = "lm",  
            aes(color = Shape))+  
  labs(x = "X",  
       y = "Y",  
       title = "An Example Plot")+  
  facet_wrap(~Shape)+  
  theme_bw()
```

An Example Plot

